

1 Emergent Misalignment and Factored Belief States

The basic empirical facts about Emergent misalignment are:

1. Fact 1: Finetuning a model on a narrowly misaligned dataset (e.g. insecure code) leads to broad misalignment across most text generations [BTW⁺25]. This works both for model that have undergone safety training, and those that have not ([BTW⁺25][Sec 4.8])
2. Fact 2: This can be controlled by a low rank operation on the weights (i.e. a rank-1 LoRA or even more simply - a steering vector) [STRN25, TST⁺25]. Recent work has also identified specific “persona features” in activation space that mediate this effect [WDITW⁺25].

These two fact seem to be precisely what we would expect if alignment was acting as an ‘almost’ factored process. In a cartoon of RLHF, models receive reward for text generation that was aligned, and penalties for unaligned text generation. If we assume that:

Assumption 1. *RLHF can be considered approximately equivalent to a supervised learning process in which a model sees the pair (capability, alignment tag) with each element being*

Then we should expect that the model will have a belief state factored into $\eta = v_c \otimes v_a$, where v_c is a vector in the capability space, and v_a is an element of the alignment space. We can then use a simple procedure to steer each factor independently:

1. Train a linear probe $\mathcal{L} = Wx + b$ from the residual stream activations x to the belief state η .
2. Perform a linear transformation S which takes the belief state η to its factorized form:
$$S\eta = v_c \otimes v_a.$$
TODO: check the below Caution: I think this is easy to do for rank-1 SVDs, but is harder for a general block diagonal matrix since this is basically the equivalent of multipartite entanglement. Hmm - that seems wrong a single cut is bi-partite entanglement.
3. Steer the alignment vector v_a .
4. Transform back into the residual stream via: $\mathcal{L}'\eta = x'$, solving the inverse problem of the original map.
5. Generate model output and compare to unmodified.

The hypothesis is then that this is what is happening in Fact 2 - the fact that the model has a factored belief state is what allows a rank-1 LoRA to steer the model to be broadly misaligned. Factored belief states, then, explain emergent misalignment. Moreover, this motivates a recipe for avoiding EM and ‘fragile’ alignment generally. In order for it to be difficult to misalign a

model, it should difficult to factor the models belief state into an ‘aligned’ and a ‘capability space’. This motivates pursuing alignment training which is more complicated than simply tagging behaviours as ‘aligned’ or ‘misaligned’. Perhaps the model personas literature [[Ant24](#), [Ant25](#)], and “Character training” [[MBLH25](#)] in particular can be thought of in this way.

1.1 Project plan

A project plan to test whether this story is true could be:

1. Test whether a model trained on a factored process, when finetuned on a fixed state of the second factor (i.e. the misaligned state) exhibits behaviour fixed to one part of the process - this seems to straightforwardly be true.
2. Test whether a factored belief state implies that each factor can be steered individually, following the procedure outlined.
3. Test whether a rank-1 LoRA, or even more simply a steering vector, acts on the factored subspace.
4. If this is true, then we can test Assumption 1 - whether RLHF can be treated in a largely similar way. To understand this, we would have to understand how factored belief states evolve under RL and see if it is largely similar.

1.2 Predictions

Is there a way of crudely testing whether this intuition holds on mid-scale (i.e Gemini-2B, Qwen-2.5B etc...) language models?

1. I think this predicts that a model which has not undergone safety training will not exhibit emergent misalignment. This is explicitly NOT what is found in the original emergent misalignment paper [[BTW⁺25](#)]. I spoke with Daniel Tan, one of the authors, about this and he suspects this is because of training data leakage - i.e. the model knows that this

2 ‘Almost factored’ processes

2.1 Dictionary between Emergent Misalignment and Almost factored belief

2.2 Partition-preserving tensor-product GHMMs (block invariance)

GHMM notation. We follow the GHMM notation of [arXiv:2507.07432](https://arxiv.org/abs/2507.07432). Let \mathcal{X} be a finite alphabet, let $(T(x))_{x \in \mathcal{X}}$ be the transfer matrices of a (finite-state) generalized hidden Markov model (GHMM), and let $\langle\!\langle \eta(\emptyset) |$ be the initial latent row vector. The net transition operator is

$$T := \sum_{x \in \mathcal{X}} T(x), \quad (2.1)$$

which has eigenvalue 1 with associated right eigenvector $|1\rangle\rangle$ satisfying $T|1\rangle\rangle = |1\rangle\rangle$. We normalize $\langle\!\langle \eta(\emptyset) |$ so that $\langle\!\langle \eta(\emptyset) | 1 \rangle\rangle = 1$.

For a word $w = x_{1:\ell}$, define $T(w) := T(x_1) \cdots T(x_\ell)$. Then

$$\Pr(X_{1:\ell} = w) = \langle\!\langle \eta(\emptyset) | T(w) | 1 \rangle\rangle. \quad (2.2)$$

The corresponding (normalized) predictive vector is

$$\langle\!\langle \eta(w) | := \frac{\langle\!\langle \eta(\emptyset) | T(w)}{\langle\!\langle \eta(\emptyset) | T(w) | 1 \rangle\rangle}. \quad (2.3)$$

In the HMM special case, the $T(x)$ are substochastic matrices with entries $T_{s,s'}^{(x)} = \Pr(s', x | s)$, and T is row-stochastic so $|1\rangle\rangle = \mathbf{1}$.

Intuition for ‘Almost factored processes’ (AFPs) For simplicity, let us consider the total state space as consisting of the tensor product of two spaces, A and B such that a general state vector lives in the tensor product state:

$$\eta = \eta_A \otimes \eta_B \quad (2.4)$$

As an intuition pump, we imagine that A and B are both “world models” of different capabilities (bad word). Previously we studied ‘Purely Factored Processes’ (PFPs). Intuitively, these processes represent sequences which contain distinct capabilities - for example, space A may represent the space of mathematical symbols, and space B the set of proper names. The key property of these processes is that knowing any piece of information about A does not tell you any information about B .

For Emergent Misalignment, however, this assumption should not hold. As another intuition pump, we take space A to represent a model persona - for simplicity consisting of an “aligned” space A_1 and a misaligned space A_2 such that $A = A_1 \oplus A_2$. We imagine space B to be the world model associated with a given capability - producing code say. We imagine we can participate that space into a space of “secure” code B_1 and insecure code B_2 . In this case we expect the intuition of factored process to break down strongly - knowing that the model is in its ‘misaligned personas’ strongly predicts that the model will also produce ‘insecure’ code.

Defining AFPs To model this dependence, we must generalize PFPs to a family of ’Almost Factored Processes’ (AFPs). In general, we consider F factor spaces $I^{(f)}$ (e.g. $I^{(1)} = A$, $I^{(2)} = B$, \dots), each of which is partitioned into N subspaces $I_i^{(f)}$ so that $I^{(f)} = \bigoplus_{i=1}^N I_i^{(f)}$. The general space of belief states V is then given by:

$$V = \bigotimes_{f=1}^F \left(\bigoplus_{i=1}^N I_i^{(f)} \right) \quad (2.5)$$

which distributes as:

$$V \cong \bigoplus_{\mathbf{i} \in \{1, \dots, N\}^F} \bigotimes_{f=1}^F I_{i_f}^{(f)}, \quad (2.6)$$

where $\mathbf{i} = (i_1, \dots, i_F)$ indexes the choice of one sector in each factor. Equivalently, we can index the sum by functions $\sigma : \{1, \dots, F\} \rightarrow \{1, \dots, N\}$:

$$V \cong \bigoplus_{\sigma : \{1, \dots, F\} \rightarrow \{1, \dots, N\}} \bigotimes_{f=1}^F I_{\sigma(f)}^{(f)}. \quad (2.7)$$

As is, this is just a relabelling of the underlying spaces. To give the partition operational significance, we impose the additional constraint that the transition operators respect the sector decomposition. Writing

$$V_\sigma := \bigotimes_{f=1}^F I_{\sigma(f)}^{(f)}, \quad V \cong \bigoplus_{\sigma : \{1, \dots, F\} \rightarrow \{1, \dots, N\}} V_\sigma, \quad (2.8)$$

we require *partition preservation*:

$$T(x)(V_\sigma) \subseteq V_\sigma, \quad \forall x \in \mathcal{X}, \forall \sigma : \{1, \dots, F\} \rightarrow \{1, \dots, N\}. \quad (2.9)$$

Equivalently, each $T(x)$ is block diagonal with respect to the direct-sum decomposition $V = \bigoplus_\sigma V_\sigma$.

Importantly, this is *not* the same as requiring full factorization: we do *not* require that a belief state in V_σ can be written as a pure tensor $\bigotimes_f v^{(f)}$, nor do we require that the restricted dynamics on a block factorizes as a tensor product of single-factor maps.

Canonical example: Two factored spaces spaces, two subspaces for each Let A and B be the hidden state spaces of two processes. A standard “factored” joint process on $A \otimes B$ has symbol-indexed operators of the form

$$T(x) = T^A(x_A) \otimes T^B(x_B), \quad (2.10)$$

where the observed symbol x encodes a pair (x_A, x_B) , and $T^A(\cdot)$, $T^B(\cdot)$ are the single-process GHMM transfer matrices.

A partition of the state spaces. Assume that each factor admits a direct-sum decomposition

$$A = A_1 \oplus A_2, \quad B = B_1 \oplus B_2. \quad (2.11)$$

By bilinearity of the tensor product, there is a canonical identification

$$A \otimes B \cong (A_1 \otimes B_1) \oplus (A_1 \otimes B_2) \oplus (A_2 \otimes B_1) \oplus (A_2 \otimes B_2). \quad (2.12)$$

We write $V_{ij} := A_i \otimes B_j$ for these four sectors.

Partition-preserving (block-invariant) dynamics. We say that a joint GHMM on $A \otimes B$ *preserves the partition* if, for every symbol $x \in \mathcal{X}$, each sector V_{ij} is invariant under $T(x)$:

$$T(x)(V_{ij}) \subseteq V_{ij}, \quad \forall x \in \mathcal{X}, \forall i, j \in \{1, 2\}. \quad (2.13)$$

Equivalently, after choosing a basis adapted to the direct sum (2.12), each $T(x)$ is block diagonal with respect to the 4-sector decomposition:

$$T(x) = \begin{pmatrix} T_{11}(x) & 0 & 0 & 0 \\ 0 & T_{12}(x) & 0 & 0 \\ 0 & 0 & T_{21}(x) & 0 \\ 0 & 0 & 0 & T_{22}(x) \end{pmatrix}, \quad T_{ij}(x) : V_{ij} \rightarrow V_{ij}. \quad (2.14)$$

We emphasize that this only constrains the *direct-sum* structure: within each block V_{ij} , the dynamics $T_{ij}(x)$ can be arbitrary and need not factorize as a tensor product of maps on A_i

and B_j . Likewise, a belief state supported on V_{ij} need not be a pure tensor in $A_i \otimes B_j$. This constraint is stronger than merely preserving a *sum* of sectors (e.g. $V_{11} \oplus V_{22}$); it enforces that the four bilinear components in the expansion

$$(a_1 + a_2) \otimes (b_1 + b_2) = a_1 \otimes b_1 + a_1 \otimes b_2 + a_2 \otimes b_1 + a_2 \otimes b_2 \quad (2.15)$$

do not mix under the dynamics.

Example: a partition-preserving Z1R-like single process. Consider a 3-state process with hidden-state ordering $(S0, S1, SR)$ and a binary alphabet $\{0, 1\}$. We impose the partition

$$\text{span}\{S0\} \oplus \text{span}\{S1, SR\}. \quad (2.16)$$

A simple Z1R-like choice that preserves this split, while retaining the emission probabilities of the usual Z1R (i.e. $\Pr(0 | S0) = 1$, $\Pr(1 | S1) = 1$, and $\Pr(0 | SR) = \Pr(1 | SR) = \frac{1}{2}$), is

$$T'_0 = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & \frac{1}{2} & 0 \end{pmatrix}, \quad T'_1 = \begin{pmatrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & \frac{1}{2} & 0 \end{pmatrix}. \quad (2.17)$$

Note that $T' := T'_0 + T'_1$ is row-stochastic:

$$T' = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix}, \quad T' \mathbf{1} = \mathbf{1}. \quad (2.18)$$

Two-process joint model with a 4-symbol alphabet. Let both factors A and B be copies of the process (2.17). The joint hidden space is $A \otimes B$ with basis ordered lexicographically:

$$(S0, S0), (S0, S1), (S0, SR), (S1, S0), (S1, S1), (S1, SR), (SR, S0), (SR, S1), (SR, SR). \quad (2.19)$$

We encode the pair of emitted bits $(x_A, x_B) \in \{0, 1\}^2$ as a single observed symbol in a 4-letter alphabet

$$\mathbf{A} = (0, 0), \quad \mathbf{B} = (0, 1), \quad \mathbf{C} = (1, 0), \quad \mathbf{D} = (1, 1). \quad (2.20)$$

The joint GHMM transfer matrices are then defined by Kronecker products:

$$T_{\mathbf{A}} = T'_0 \otimes T'_0, \quad T_{\mathbf{B}} = T'_0 \otimes T'_1, \quad T_{\mathbf{C}} = T'_1 \otimes T'_0, \quad T_{\mathbf{D}} = T'_1 \otimes T'_1. \quad (2.21)$$

By construction, each T_A, T_B, T_C, T_D is entrywise nonnegative and the net transition operator

$$T_{\text{tot}} = T_A + T_B + T_C + T_D = (T'_0 + T'_1) \otimes (T'_0 + T'_1) = T' \otimes T' \quad (2.22)$$

satisfies $T_{\text{tot}} \mathbf{1} = \mathbf{1}$.

Partition preservation in the joint model. Under the joint partition induced by $\text{span}\{S0\} \oplus \text{span}\{S1, SR\}$ in each factor,

$$A \otimes B = (\text{span}\{S0\} \otimes \text{span}\{S0\}) \oplus (\text{span}\{S0\} \otimes \text{span}\{S1, SR\}) \oplus (\text{span}\{S1, SR\} \otimes \text{span}\{S0\}) \oplus (\text{span}\{S1, SR\} \otimes \text{span}\{S1, SR\}) \quad (2.23)$$

each symbol operator in (2.21) leaves all four sectors invariant, hence is block diagonal in an adapted basis as in (2.14). This gives an explicit family of partition-preserving joint GHMMs that reduce to a standard Kronecker-product construction within each block.

2.3 Density matrix formulation

For the two-factor HMM case, it is sometimes helpful to represent a belief state as a (classical) density matrix on $A \otimes B$. Fix the computational basis $\{|s_A\rangle\}_{s_A \in \{S0, S1, SR\}}$ for A and $\{|s_B\rangle\}_{s_B \in \{S0, S1, SR\}}$ for B , and write $|s_A, s_B\rangle := |s_A\rangle \otimes |s_B\rangle$.

Classical density matrix and correlations. Given a history w , the predictive state induces a joint distribution $p_w(s_A, s_B) = \Pr(S_A = s_A, S_B = s_B \mid w)$ on hidden states. The corresponding density matrix is diagonal:

$$\rho_{AB}(w) := \sum_{s_A, s_B} p_w(s_A, s_B) |s_A, s_B\rangle \langle s_A, s_B|. \quad (2.24)$$

In the lexicographic basis used above,

$$\rho_{AB}(w) = \text{diag}\left(p_{00}, p_{01}, p_{0R}, p_{10}, p_{11}, p_{1R}, p_{R0}, p_{R1}, p_{RR}\right), \quad (2.25)$$

where, e.g., $p_{0R} := p_w(S0, SR)$ and $\sum p_{ab} = 1$. The marginals are $\rho_A = \text{Tr}_B \rho_{AB}$ and $\rho_B = \text{Tr}_A \rho_{AB}$. The factors are *uncorrelated* iff $\rho_{AB} = \rho_A \otimes \rho_B$ (equivalently $p_w(s_A, s_B) = p_w(s_A) p_w(s_B)$). Otherwise, the correlation is purely classical: it appears in the diagonal entries (joint probabilities), not in off-diagonal coherences.

Block structure from the partition. With the two-subspace partition $\text{span}\{S0\} \oplus \text{span}\{S1, SR\}$ in each factor, $\rho_{AB}(w)$ is block diagonal with block sizes 1, 2, 2, 4 corresponding to the sectors

in (2.12). The total weight in each block is the coarse-grained joint distribution over partition labels.

Almost-factored $\mathbf{Z1R} \times \mathbf{Z1R}$ belief state. An “almost factored” (but still classical/diagonal) belief state can have strong correlations between the partitions. For example, putting all mass on the “matched” sectors V_{11} and V_{22} gives

$$\rho_{AB}^{\text{AFP}} = p |S0, S0\rangle \langle S0, S0| + \frac{1-p}{4} \sum_{s \in \{S1, SR\}} \sum_{t \in \{S1, SR\}} |s, t\rangle \langle s, t|, \quad (2.26)$$

which is block diagonal (so it respects the partition) but generally satisfies $\rho_{AB}^{\text{AFP}} \neq \rho_A \otimes \rho_B$, hence encodes classical correlation between the factors.

3 Reproducing emergent misalignment fine-tuning

References

- [Ant24] Anthropic. The claude model spec: Character training. <https://www.anthropic.com/research/clause-character>, 2024.
- [Ant25] Anthropic. Persona vectors: Monitoring and controlling character traits. <https://www.anthropic.com/research/persona-vectors>, 2025.
- [BTW⁺25] Jan Betley, Daniel Tan, Niels Warncke, Anna Sztyber-Betley, Xuchan Bao, Martin Soto, Nathan Labenz, and Owain Evans. Emergent misalignment: Narrow finetuning can produce broadly misaligned llms. *arXiv preprint arXiv:2502.17424*, 2025.
- [MBLH25] Sharan Maiya, Henning Bartsch, Nathan Lambert, and Evan Hubinger. Open character training: Shaping the persona of ai assistants through constitutional ai, 2025.
- [STRN25] Anna Soligo, Liv Turner, Senthooran Rajamanoharan, and Neel Nanda. Convergent linear representations of emergent misalignment. *arXiv preprint arXiv:2506.11618*, 2025.
- [TST⁺25] Liv Turner, Anna Soligo, Jessica Taylor, Senthooran Rajamanoharan, and Neel Nanda. Model organisms for emergent misalignment. *arXiv preprint arXiv:2506.11613*, 2025.

- [WDLTW⁺25] Miles Wang, Tom Dupré la Tour, Olivia Watkins, Alex Makelov, Ryan A Chi, Samuel Miserendino, Johannes Heidecke, Tejal Patwardhan, and Dan Mossing. Persona features control emergent misalignment. *arXiv preprint arXiv:2506.19823*, 2025.